

Notebook Gemini



```
import pandas as pd

file_path = '/content/Twitter_Data.csv'
df = pd.read_csv(file_path, on_bad_lines='skip')
display(df.head())
```

	clean_text	category	
0	when modi promised "minimum government maximum...	-1.0	
1	talk all the nonsense and continue all the dra...	0.0	
2	what did just say vote for modi welcome bjp t...	1.0	
3	asking his supporters prefix chowkidar their n...	1.0	
4	answer who among these the most powerful world...	1.0	

```
import re
import numpy as np

# Function to remove URLs
def remove_urls(text):
    if pd.isna(text):
        return ""
    url_pattern = re.compile(r'https?://\S+|www\.\S+')
    return url_pattern.sub(r'', str(text))

# Function to remove mentions
def remove_mentions(text):
    if pd.isna(text):
        return ""
    mention_pattern = re.compile(r'@\w+')
    return mention_pattern.sub(r'', str(text))

# Apply the preprocessing steps to the 'clean_text' column
df['processed_text'] = df['clean_text'].apply(remove_urls)
df['processed_text'] = df['processed_text'].apply(remove_mentions)

# Display the DataFrame with the new processed_text column
display(df[['clean_text', 'processed_text']].head())
```

	clean_text	processed_text	
0	when modi promised "minimum government maximum..."	when modi promised "minimum government maximum..."	
1	talk all the nonsense and continue all the dra...	talk all the nonsense and continue all the dra...	
2	what did just say vote for modi welcome bjp t...	what did just say vote for modi welcome bjp t...	
3	asking his supporters prefix chowkidar their n...	asking his supporters prefix chowkidar their n...	
4	answer who among these the most powerful world...	answer who among these the most powerful world...	

```
import nltk

# Download necessary NLTK data
nltk.download('punkt', quiet=True, raise_on_error=False)
nltk.download('averaged_perceptron_tagger_eng', quiet=True, raise_on_error=False)
nltk.download('punkt_tab', quiet=True, raise_on_error=False)

True
```

```
# Function to perform POS tagging
def pos_tag_text(text):
    if pd.isna(text) or text == "":
        return []
    tokens = nltk.word_tokenize(text)
    return nltk.pos_tag(tokens)

# Apply POS tagging to the 'processed_text' column
df['pos_tagged_text'] = df['processed_text'].apply(pos_tag_text)
```

```
# Display the DataFrame with the new pos_tagged_text column
display(df[['processed_text', 'pos_tagged_text']].head())
```

	processed_text	pos_tagged_text
0	when modi promised "minimum government maximum..."	[(when, WRB), (modi, NN), (promised, VBD), ("...,
1	talk all the nonsense and continue all the dra...	[(talk, NN), (all, PDT), (the, DT), (nonsense,...
2	what did just say vote for modi welcome bjp t...	[(what, WP), (did, VBD), (just, RB), (say, VB)...
3	asking his supporters prefix chowkidar their n...	[(asking, VBG), (his, PRP\$), (supporters, NNS)...
4	answer who among these the most powerful world...	[(answer, NN), (who, WP), (among, IN), (these,...

```
from collections import defaultdict

# Initialize dictionaries for HMM parameters
initial_probabilities = defaultdict(lambda: 0)
transition_probabilities = defaultdict(lambda: defaultdict(lambda: 0))
emission_probabilities = defaultdict(lambda: defaultdict(lambda: 0))
tag_counts = defaultdict(lambda: 0)
vocabulary = set()

# Iterate through each POS-tagged tweet
for tagged_tweet in df['pos_tagged_text']:
    if not tagged_tweet:
        continue

    # Increment initial probability for the first tag in the tweet
    first_word, first_tag = tagged_tweet[0]
    initial_probabilities[first_tag] += 1
    tag_counts[first_tag] += 1
    vocabulary.add(first_word.lower())

    # Iterate through the rest of the words for transition and emission probabilities
    for i in range(len(tagged_tweet)):
        word, tag = tagged_tweet[i]
        vocabulary.add(word.lower())

        emission_probabilities[tag][word.lower()] += 1

        if i > 0:
            prev_word, prev_tag = tagged_tweet[i-1]
            transition_probabilities[prev_tag][tag] += 1
            tag_counts[tag] += 1 # Count all tags for normalization

    # Normalize probabilities

    # Normalize initial probabilities
    total_initial_tags = sum(initial_probabilities.values())
    if total_initial_tags > 0:
        for tag in initial_probabilities:
            initial_probabilities[tag] /= total_initial_tags

    # Normalize transition probabilities
    for prev_tag, next_tag_counts in transition_probabilities.items():
        total_transitions = sum(next_tag_counts.values())
        if total_transitions > 0:
            for next_tag in next_tag_counts:
                transition_probabilities[prev_tag][next_tag] /= total_transitions

    # Normalize emission probabilities
    for tag, word_counts in emission_probabilities.items():
        total_emissions = sum(word_counts.values())
        if total_emissions > 0:
            for word in word_counts:
                emission_probabilities[tag][word] /= total_emissions

print("HMM Parameters Built Successfully!")
print("\nTop 5 Initial Probabilities:")
for tag, prob in sorted(initial_probabilities.items(), key=lambda item: item[1], reverse=True)[:5]:
    print(f" {tag}: {prob:.4f}")

print("\nTop 5 Transition Probabilities (e.g., from NN):")
if 'NN' in transition_probabilities:
```

```

        for next_tag, prob in sorted(transition_probabilities['NN'].items(), key=lambda item: item[1], reverse=True)[:5]:
            print(f"  NN -> {next_tag}: {prob:.4f}")
    else:
        print("  'NN' tag not found in transitions.")

print("\nTop 5 Emission Probabilities (e.g., for NN):")
if 'NN' in emission_probabilities:
    for word, prob in sorted(emission_probabilities['NN'].items(), key=lambda item: item[1], reverse=True)[:5]:
        print(f"  NN -> {word}: {prob:.4f}")
else:
    print("  'NN' tag not found in emissions.")

```

HMM Parameters Built Successfully!

Top 5 Initial Probabilities:

NN: 0.3298
JJ: 0.1157
RB: 0.1023
DT: 0.0759
NNS: 0.0688

Top 5 Transition Probabilities (e.g., from NN):

NN -> NN: 0.3551
NN -> IN: 0.0930
NN -> NNS: 0.0617
NN -> VBD: 0.0568
NN -> RB: 0.0551

Top 5 Emission Probabilities (e.g., for NN):

NN -> modi: 0.1005
NN -> india: 0.0140
NN -> congress: 0.0119
NN -> bjp: 0.0112
NN -> money: 0.0096

▼ Inspect Transition Probabilities for a specific previous tag

```

#@title Inspect Transition Probabilities for a specific prev

# You can change the 'prev_tag_to_inspect' to any valid POS
prev_tag_to_inspect = 'NN' # @param {type:"string"}

print(f"\nTop 10 Transition Probabilities FROM '{prev_tag_to_inspect}'")
if prev_tag_to_inspect in transition_probabilities:
    sorted_transitions = sorted(
        transition_probabilities[prev_tag_to_inspect].items(
            key=lambda item: item[1],
            reverse=True
        )[:10]

    if sorted_transitions:
        for next_tag, prob in sorted_transitions:
            print(f"  {prev_tag_to_inspect} -> {next_tag}: {prob:.4f}")
    else:
        print(f"  No observed transitions from '{prev_tag_to_inspect}'")
else:
    print(f"  '{prev_tag_to_inspect}' tag not found in trans")

```

prev_tag_to_inspect

NN

[Hide code](#)

Top 10 Transition Probabilities FROM 'NN':

NN -> NN: 0.3551
NN -> IN: 0.0930
NN -> NNS: 0.0617
NN -> VBD: 0.0568
NN -> RB: 0.0551
NN -> JJ: 0.0537
NN -> CC: 0.0520
NN -> VBZ: 0.0461
NN -> MD: 0.0314
NN -> DT: 0.0292

```

from collections import Counter
import itertools

# Flatten the list of tokenized words from 'pos_tagged_text'
# We only need the word, not the tag, for overall word frequency

```

```

all_words = list(itertools.chain.from_iterable(
    [word for word, tag in tweet_tags] for tweet_tags in df['pos_tagged_text'] if tweet_tags
))

# Convert all words to lowercase for consistent counting
all_words_lower = [word.lower() for word in all_words]

# Count word frequencies
word_counts = Counter(all_words_lower)

print("Total unique words in vocabulary:", len(word_counts))
print("Total words (including duplicates):", len(all_words_lower))

print("\nTop 20 most common words:")
for word, count in word_counts.most_common(20):
    print(f" '{word}': {count}")

print("\nTop 20 rarest words (appearing only once):")
rare_words = [(word, count) for word, count in word_counts.items() if count == 1]
if rare_words:
    for word, count in sorted(rare_words, key=lambda item: item[0])[::20]: # Sort alphabetically for consistent output
        print(f" '{word}': {count}")
else:
    print(" No words appear only once.")

# We can also inspect the words with the lowest emission probabilities for a specific tag
# For example, words rarely emitted by 'NN'
print("\nTop 10 words with lowest emission probability for 'NN' (if NN is present and has emissions):")
if 'NN' in emission_probabilities and emission_probabilities['NN']:
    sorted_nn_emissions = sorted(
        emission_probabilities['NN'].items(),
        key=lambda item: item[1]
    )[::10]
    for word, prob in sorted_nn_emissions:
        print(f" 'NN' -> '{word}': {prob:.6f}")
else:
    print(" 'NN' tag not found or has no emission probabilities.")

```

Total unique words in vocabulary: 33517
Total words (including duplicates): 499980

Top 20 most common words:

```
'modi': 22863
'the': 15405
'and': 10159
'for': 7966
'you': 6336
'will': 4897
'not': 4681
'this': 4540
'that': 4368
'are': 4270
'india': 3800
'has': 3170
'with': 3130
'have': 3084
'but': 2815
'from': 2700
'all': 2680
'''': 2627
'bjp': 2492
'his': 2452
```

Top 20 rarest words (appearing only once):

```
'000waiting': 1  
'025': 1  
'0420462818': 1  
'04280': 1  
'05852234246': 1  
'05yrs': 1  
'062': 1  
'081': 1  
'08th': 1  
'0deposit': 1  
'100000': 1  
'100000₹': 1  
'1000rs': 1  
'100100': 1  
'100agree': 1  
'100because': 1
```

```

100cr : 1
'100dayscan': 1
'100it': 1
'100paise': 1

```

```

Top 10 words with lowest emission probability for 'NN' (if NN is present and has emissions):
NN -> 'constituency2': 0.000007
NN -> 'tuthukudi': 0.000007
NN -> 'thuthukudi': 0.000007
NN -> 'leadershipwho': 0.000007
NN -> 'khammam': 0.000007
NN -> 'condidate': 0.000007
NN -> 'sonbadra': 0.000007
NN -> 'modiganga': 0.000007
NN -> 'rejuvenation': 0.000007
NN -> 'repressive': 0.000007

```

```

def viterbi(words, initial_probs, transition_probs, emission_probs, all_tags):
    # Initialize Viterbi path and probabilities
    # Viterbi_path[tag][t] stores the tag sequence ending with 'tag' at time 't'
    # Viterbi_prob[tag][t] stores the probability of the most likely path ending with 'tag' at time 't'
    T = len(words)
    Viterbi_prob = [{tag: 0.0 for tag in all_tags} for _ in range(T)]
    Viterbi_path = [{tag: [] for tag in all_tags} for _ in range(T)]

    # Step 1: Initialization for t = 0
    for tag in all_tags:
        # Initial probability of starting with 'tag'
        init_p = initial_probs.get(tag, 1e-10) # Use a small non-zero value for unseen tags

        # Emission probability of the first word given the tag
        word_lower = words[0].lower()
        emit_p = emission_probs.get(tag, {}).get(word_lower, 1e-10) # Use a small non-zero value for unseen words/tags

        Viterbi_prob[0][tag] = init_p * emit_p
        Viterbi_path[0][tag] = [tag]

    # Step 2: Recursion for t = 1 to T-1
    for t in range(1, T):
        word_lower = words[t].lower()
        for current_tag in all_tags:
            max_prob = 0.0
            best_prev_tag = None

            # Emission probability of the current word given the current tag
            emit_p = emission_probs.get(current_tag, {}).get(word_lower, 1e-10)

            for prev_tag in all_tags:
                # Probability of path ending with 'prev_tag' at t-1
                prob_prev_path = Viterbi_prob[t-1][prev_tag]

                # Transition probability from 'prev_tag' to 'current_tag'
                trans_p = transition_probs.get(prev_tag, {}).get(current_tag, 1e-10)

                # Calculate total probability for this path
                current_path_prob = prob_prev_path * trans_p * emit_p

                if current_path_prob > max_prob:
                    max_prob = current_path_prob
                    best_prev_tag = prev_tag

            Viterbi_prob[t][current_tag] = max_prob
            if best_prev_tag:
                Viterbi_path[t][current_tag] = Viterbi_path[t-1][best_prev_tag] + [current_tag]
            else:
                Viterbi_path[t][current_tag] = [current_tag] # Fallback if no prev tag found (shouldn't happen with 1e-10)

    # Step 3: Termination
    max_prob = 0.0
    best_last_tag = None

    for tag in all_tags:
        if Viterbi_prob[T-1][tag] > max_prob:
            max_prob = Viterbi_prob[T-1][tag]
            best_last_tag = tag

    if best_last_tag:
        return Viterbi_path[T-1][best_last_tag]

```

```

        else:
            return [] # Should not happen if words is not empty

# Get all unique tags from the training data
all_tags = set()
for tagged_tweet in df['pos_tagged_text']:
    for word, tag in tagged_tweet:
        all_tags.add(tag)

print("Viterbi algorithm defined and ready for use!")

```

Viterbi algorithm defined and ready for use!

Snapshot of HMM Parameters

Below are ways to inspect the `initial_probabilities`, `transition_probabilities`, and `emission_probabilities` that have been computed.

```

print("\n--- Initial Probabilities (Top 10) ---")
# Sorting to see the most probable initial tags
for tag, prob in sorted(initial_probabilities.items(), key=lambda item: item[1], reverse=True)[:10]:
    print(f" Tag: {tag}, Probability: {prob:.4f}")

print("\n--- Transition Probabilities (Examples from 'NN', 'VB', 'JJ') ---")
# Inspecting a few example tags for their transitions
for prev_tag_example in ['NN', 'VB', 'JJ']:
    if prev_tag_example in transition_probabilities:
        print(f" Top 5 transitions from '{prev_tag_example}':")
        sorted_transitions = sorted(
            transition_probabilities[prev_tag_example].items(),
            key=lambda item: item[1],
            reverse=True
        )[:5]
        if sorted_transitions:
            for next_tag, prob in sorted_transitions:
                print(f"    {prev_tag_example} -> {next_tag}: {prob:.4f}")
        else:
            print(f"    No observed transitions from '{prev_tag_example}'.")

    else:
        print(f"    '{prev_tag_example}' tag not found in transition probabilities.")

print("\n--- Emission Probabilities (Examples for 'NN', 'VB', 'JJ') ---")
# Inspecting a few example tags for their emissions
for tag_example in ['NN', 'VB', 'JJ']:
    if tag_example in emission_probabilities:
        print(f" Top 5 emitted words for '{tag_example}':")
        sorted_emissions = sorted(
            emission_probabilities[tag_example].items(),
            key=lambda item: item[1],
            reverse=True
        )[:5]
        if sorted_emissions:
            for word, prob in sorted_emissions:
                print(f"    {tag_example} -> '{word}': {prob:.4f}")
        else:
            print(f"    No observed emissions for '{tag_example}'.")

    else:
        print(f"    '{tag_example}' tag not found in emission probabilities.")

```

```

--- Initial Probabilities (Top 10) ---
Tag: NN, Probability: 0.3298
Tag: JJ, Probability: 0.1157
Tag: RB, Probability: 0.1023
Tag: DT, Probability: 0.0759
Tag: NNS, Probability: 0.0688
Tag: VB, Probability: 0.0461
Tag: IN, Probability: 0.0379
Tag: WRB, Probability: 0.0345
Tag: PRP, Probability: 0.0302
Tag: MD, Probability: 0.0202

```

```

--- Transition Probabilities (Examples from 'NN', 'VB', 'JJ') ---

```

```

Top 5 transitions from 'NN':
NN -> NN: 0.3551
NN -> IN: 0.0930
NN -> NNS: 0.0617
NN -> VBD: 0.0568
NN -> RB: 0.0551
Top 5 transitions from 'VB':
VB -> NN: 0.1897
VB -> DT: 0.1501
VB -> JJ: 0.1174
VB -> IN: 0.0818
VB -> PRP$: 0.0807
Top 5 transitions from 'JJ':
JJ -> NN: 0.5571
JJ -> NNS: 0.1561
JJ -> JJ: 0.0989
JJ -> IN: 0.0334
JJ -> RB: 0.0214

--- Emission Probabilities (Examples for 'NN', 'VB', 'JJ') ---
Top 5 emitted words for 'NN':
NN -> 'modi': 0.1005
NN -> 'india': 0.0140
NN -> 'congress': 0.0119
NN -> 'bjp': 0.0112
NN -> 'money': 0.0096
Top 5 emitted words for 'VB':
VB -> 'have': 0.0401
VB -> 'modi': 0.0370
VB -> 'get': 0.0363
VB -> 'give': 0.0299
VB -> 'make': 0.0191
Top 5 emitted words for 'JJ':
JJ -> 'modi': 0.0494
JJ -> 'poor': 0.0224
JJ -> 'indian': 0.0164
JJ -> 'narendra': 0.0142
JJ -> 'good': 0.0136

```

▼ Apply Viterbi Decoding to a Sample Tweet

Now, let's take a sample tweet, tokenize it, and then apply our `viterbi` function to find its most likely POS tag sequence.

```

# Choose a sample tweet (you can modify this tweet or select another from df['processed_text'])
sample_tweet = "modi will make india great again"

# Tokenize the sample tweet
sample_words = nltk.word_tokenize(sample_tweet)
print(f"Sample Tweet: {sample_tweet}")
print(f"Tokenized Words: {sample_words}")

# Apply Viterbi decoding
predicted_tags = viterbi(sample_words, initial_probabilities, transition_probabilities, emission_probabilities, all_tags)

# Combine words and predicted tags for display
result = list(zip(sample_words, predicted_tags))

print(f"\nPredicted POS Tag Sequence: {result}")

Sample Tweet: 'modi will make india great again'
Tokenized Words: ['modi', 'will', 'make', 'india', 'great', 'again']

Predicted POS Tag Sequence: [('modi', 'NN'), ('will', 'MD'), ('make', 'VB'), ('india', 'JJ'), ('great', 'JJ'), ('again', 'RB')]

```

Start coding or [generate](#) with AI.

